

18/09/2025

PROYECTO **MATCHAI**

A01710550 - Maxime Vilcocq Parra

A01710791 - Galo Alejandro del Rio Viggiano

A01369687 - Ana Karen Toscano Díaz

A01710367 - José Antonio López Saldaña

EQUIPO DE TRABAJO



Galo Alejandro del
Rio Viggiano

A01710791




Maxime Vilcocq
Parra

A01710550



José Antonio López
Saldaña



A01710367



Ana Karen (Max)
Toscano

A01369687

Sobre el Proyecto



CONTEXTO

Sobre el Proyecto

NUESTRA SOLUCIÓN




OBJETIVOS



Detección



Precisión



Robustez



**CÓDIGO
IMPLEMENTADO**

1. CARGA Y PREPROCESAMIENTO DE DATOS



Carga de imágenes TAC y máscaras desde archivos .npy



Normalización con windowing tipo CT



Clipping intensidades (-1500 a 500 HU)

```
# Load data
```

```
prefix = "/kaggle/input/covid-segmentation/"
```

```
images_medseg = np.load(prefix + "images_medseg.npy").astype(np.float32)
```

```
masks_medseg = np.load(prefix + "masks_medseg.npy").astype(np.int16)
```

```
NUM_CLASSES = 4
```

```
# Preprocess
```

```
def preprocess(images: np.ndarray) -> np.ndarray:
```

```
    images = np.clip(images, -1500, 500)
```

```
    mean, std = images.mean(), images.std()
```

```
    return (images - mean) / (std + 1e-8)
```

```
images_medseg = preprocess(images_medseg)
```


2. DATA SET Y DATA LOADER



Creación de la clase LungDataset

- Convierte imágenes a tensores
- Asegura máscaras indexadas (H,W)



Separación en entrenamiento y validación



Uso de WeightedRandomSampler para balancear clases minoritarias

Dataset

```
class LungDataset(Dataset):
    def __init__(self, images, masks, aug=None):
        self.images, self.masks, self.aug = images, masks, aug
        self.to_tensor = T.ToTensor() # HWC -> CHW, preserva float32

    def __len__(self):
        return len(self.images)

    def __getitem__(self, i):
        img, mask = self.images[i], self.masks[i]

        if self.aug: # aplicar augmentations si están definidos
            sample = self.aug(image=img, mask=mask)
            img, mask = sample["image"], sample["mask"]

        if img.ndim == 2: # si no hay canal explícito añadir (H, W, 1)
            img = img[..., None]

        x = self.to_tensor(img)

        if mask.ndim == 3 and mask.shape[-1] > 1: # si está one-hot pasar a índices
            mask = np.argmax(mask, axis=-1)
        y = torch.tensor(mask, dtype=torch.long)
        return x, y
```

```
# Train/Val split
```

```
n_total = len(images_medseg)
```

```
n_val = int(0.20 * n_total)
```

```
idxs = np.arange(n_total); np.random.shuffle(idxs)
```

```
train_idx, val_idx = idxs[n_val:], idxs[:n_val]
```

```
train_ds = LungDataset(images_medseg[train_idx], masks_medseg[train_idx], train_aug)
```

```
val_ds = LungDataset(images_medseg[val_idx], masks_medseg[val_idx], val_aug)
```

```
# Weighted sampler
```

```
def mask_has_lesion(mask):  
    if mask.ndim == 3 and mask.shape[-1] > 1:  
        mask = np.argmax(mask, axis=-1)  
    return int(np.any(mask > 0))
```

```
minor_presence = np.array([mask_has_lesion(m) for m in masks_medseg[train_idx]])  
sample_weights_np = np.where(minor_presence == 1, 3.0, 1.0).astype(np.float32)  
sample_weights = torch.as_tensor(sample_weights_np, dtype=torch.double)  
sampler = WeightedRandomSampler(sample_weights, num_samples=len(sample_weights), replacement=True)
```

```
BATCH = 6
```

```
pin = torch.cuda.is_available()
```

```
train_dl = DataLoader(train_ds, batch_size=BATCH, sampler=sampler, num_workers=2, pin_memory=pin)
```

```
val_dl = DataLoader(val_ds, batch_size=BATCH, shuffle=False, num_workers=2, pin_memory=pin)
```

3 . A U G M E N T A T I O N S



Entrenamiento (más variabilidad):

- Resize a 320×320
- Flips (horizontal, vertical)
- Rotaciones aleatorias (90°)
- Afine (escala, rotación, shear, traslación)
- Brillo y contraste aleatorio




Validación (solo resize para consistencia)

```
# Augmentations
```

```
IMG_SIZE = 384
```

```
train_aug = A.Compose([  
    A.Resize(IMG_SIZE, IMG_SIZE),  
    A.HorizontalFlip(p=0.5),  
    A.VerticalFlip(p=0.3),  
    A.RandomRotate90(p=0.3),  
    A.Affine(scale=(0.95,1.05), rotate=(-10,10), shear=(-5,5),  
            translate_percent=(0.0,0.03), p=0.4),  
    A.RandomBrightnessContrast(brightness_limit=0.10, contrast_limit=0.10, p=0.2),  
)  
val_aug = A.Compose([A.Resize(IMG_SIZE, IMG_SIZE)])
```

4 . M O D E L O




Arquitectura UNet++



Encoder: EfficientNet-B3 preentrenado en ImageNet



Entrada: 1 canal (grayscale)



Salida: 4 clases (fondo + 2 lesiones)

```
# Model
```

```
device = "cuda" if torch.cuda.is_available() else "cpu"  
model = smp.UnetPlusPlus(  
    encoder_name="timm-efficientnet-b3",  
    encoder_weights="imagenet",  
    in_channels=1,  
    classes=NUM_CLASSES,  
) .to(device)
```


5. FUNCIONES DE PÉRDIDA

π

CrossEntropy con pesos balanceados.

π

Dice Loss para segmentación.

π

Focal Loss para énfasis en clases difíciles ($\gamma = 2.5$).

π

Ponderación: $0.25 \cdot \text{CE} + 0.5 \cdot \text{Dice} + 0.25 \cdot \text{Focal}$

```
# Definir combinación de pérdidas = CrossEntropy ponderada + Dice + Focal
```

```
ce_loss = nn.CrossEntropyLoss(weight=ce_weights)
```

```
dice_loss = smp.losses.DiceLoss(mode="multiclass")
```

```
focal_loss = smp.losses.FocalLoss(mode="multiclass", gamma=2.5)
```

```
def criterion(y_pred, y_true, a=0.25, b=0.5, c=0.25):
```

```
    return a*ce_loss(y_pred, y_true) + b*dice_loss(y_pred, y_true) + c*focal_loss(y_pred, y_true)
```

6. OPTIMIZACIÓN Y MÉTRICAS



Optimizador: AdamW



Scheduler: OneCycleLR (ajuste dinámico del LR)



AMP (Mixed Precision) para acelerar entrenamiento



Métrica principal: Coeficiente Dice por clase

```
# Optimizer & Scheduler
```

```
opt = torch.optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
EPOCHS = 60
scheduler = torch.optim.lr_scheduler.OneCycleLR(
    opt, max_lr=5e-4, steps_per_epoch=len(train_dl), epochs=EPOCHS
)
```

```
# Configuración para entrenamiento mixto (AMP) con nueva API.
```

```
use_amp = (device == "cuda")
```

```
scaler = torch.amp.GradScaler('cuda') if use_amp else None
```

```
# Metrics
```

```
@torch.no_grad()
```

```
def dice_per_class(logits, target, num_classes=NUM_CLASSES):
```

```
    pred = logits.argmax(1)
```

```
    dices = []
```

```
    for c in range(num_classes):
```

```
        p = (pred == c).float()
```

```
        t = (target == c).float()
```

```
        inter = (p*t).sum()
```

```
        denom = p.sum() + t.sum()
```

```
        dices.append(1.0 if denom == 0 else (2*inter/denom).item())
```

```
    return dices
```

7. LOOP DE ENTRENAMIENTO



Forward y backward pass con AMP



Validación al final de cada epoch



Guardado del mejor modelo (checkpoint) según Dice promedio

```
# Training loop
```

```
best_mean_dice = -1.0
```

```
ckpt_path = "best_unetpp_b3_384.pth"
```

```
for epoch in range(1, EPOCHS+1):
```

```
    model.train()
```

```
    running = 0.0
```

```
    for x, y in train_dl:
```

```
        x, y = x.to(device, non_blocking=True), y.to(device, non_blocking=True)
```

```
        opt.zero_grad(set_to_none=True)
```

```
        if use_amp:
```

```
            with torch.amp.autocast('cuda'):
```

```
                out = model(x)
```

```
                loss = criterion(out, y)
```

```
                scaler.scale(loss).backward()
```

```
                scaler.step(opt)
```

```
                scaler.update()
```

```
        else:
```

```
            out = model(x)
```

```
            loss = criterion(out, y)
```

```
            loss.backward()
```

```
            opt.step()
```

```
        scheduler.step()
```

```
        running += loss.item()
```


```
    train_loss = running / max(len(train_dl), 1)
```

```
model.eval()
vloss, vdices = 0.0, []
autocast_ctx = torch.amp.autocast('cuda') if use_amp else contextlib.nullcontext()
with torch.no_grad(), autocast_ctx:
    for x, y in val_dl:
        x, y = x.to(device), y.to(device)
        out = model(x)
        loss = criterion(out, y)
        vloss += loss.item()
        vdices.append(dice_per_class(out, y))
vloss /= max(len(val_dl), 1)
vdices = np.mean(vdices, axis=0)
mean_dice = float(np.mean(vdices))


if mean_dice > best_mean_dice:
    best_mean_dice = mean_dice
    torch.save(model.state_dict(), ckpt_path)

print(f"Epoch {epoch:3d}/{EPOCHS} | Train {train_loss:.3f} | Val {vloss:.3f} | "
      f"Dice {np.round(vdices, 4)} | Mean {mean_dice:.4f} | Best {best_mean_dice:.4f}")
```

8. VISUALIZACIÓN DE RESULTADOS



Función visualize_batch



Muestra:

- Imagen TAC
- Máscara ground truth
- Predicción del modelo

Visualization

```
model.load_state_dict(torch.load(ckpt_path, map_location=device))  
model.eval()
```

```
@torch.no_grad()
```

```
def visualize_batch(x_cpu, y_cpu, pred_cpu, max_images=6):
```

```
    n = min(max_images, x_cpu.size(0))
```

```
    for i in range(n):
```

```
        plt.figure(figsize=(12,4))
```

```
        plt.subplot(1,3,1); plt.imshow(x_cpu[i,0], cmap="gray"); plt.title("Image"); plt.axis("off")
```

```
        plt.subplot(1,3,2); plt.imshow(y_cpu[i], cmap="jet", alpha=0.7); plt.title("Mask GT"); plt.axis("off")
```

```
        plt.subplot(1,3,3); plt.imshow(pred_cpu[i], cmap="jet", alpha=0.7); plt.title("Prediction"); plt.axis("off")
```

```
    plt.show()
```

```
with torch.no_grad():
```

```
    for b_idx, (x, y) in enumerate(val_dl):
```

```
        x = x.to(device)
```

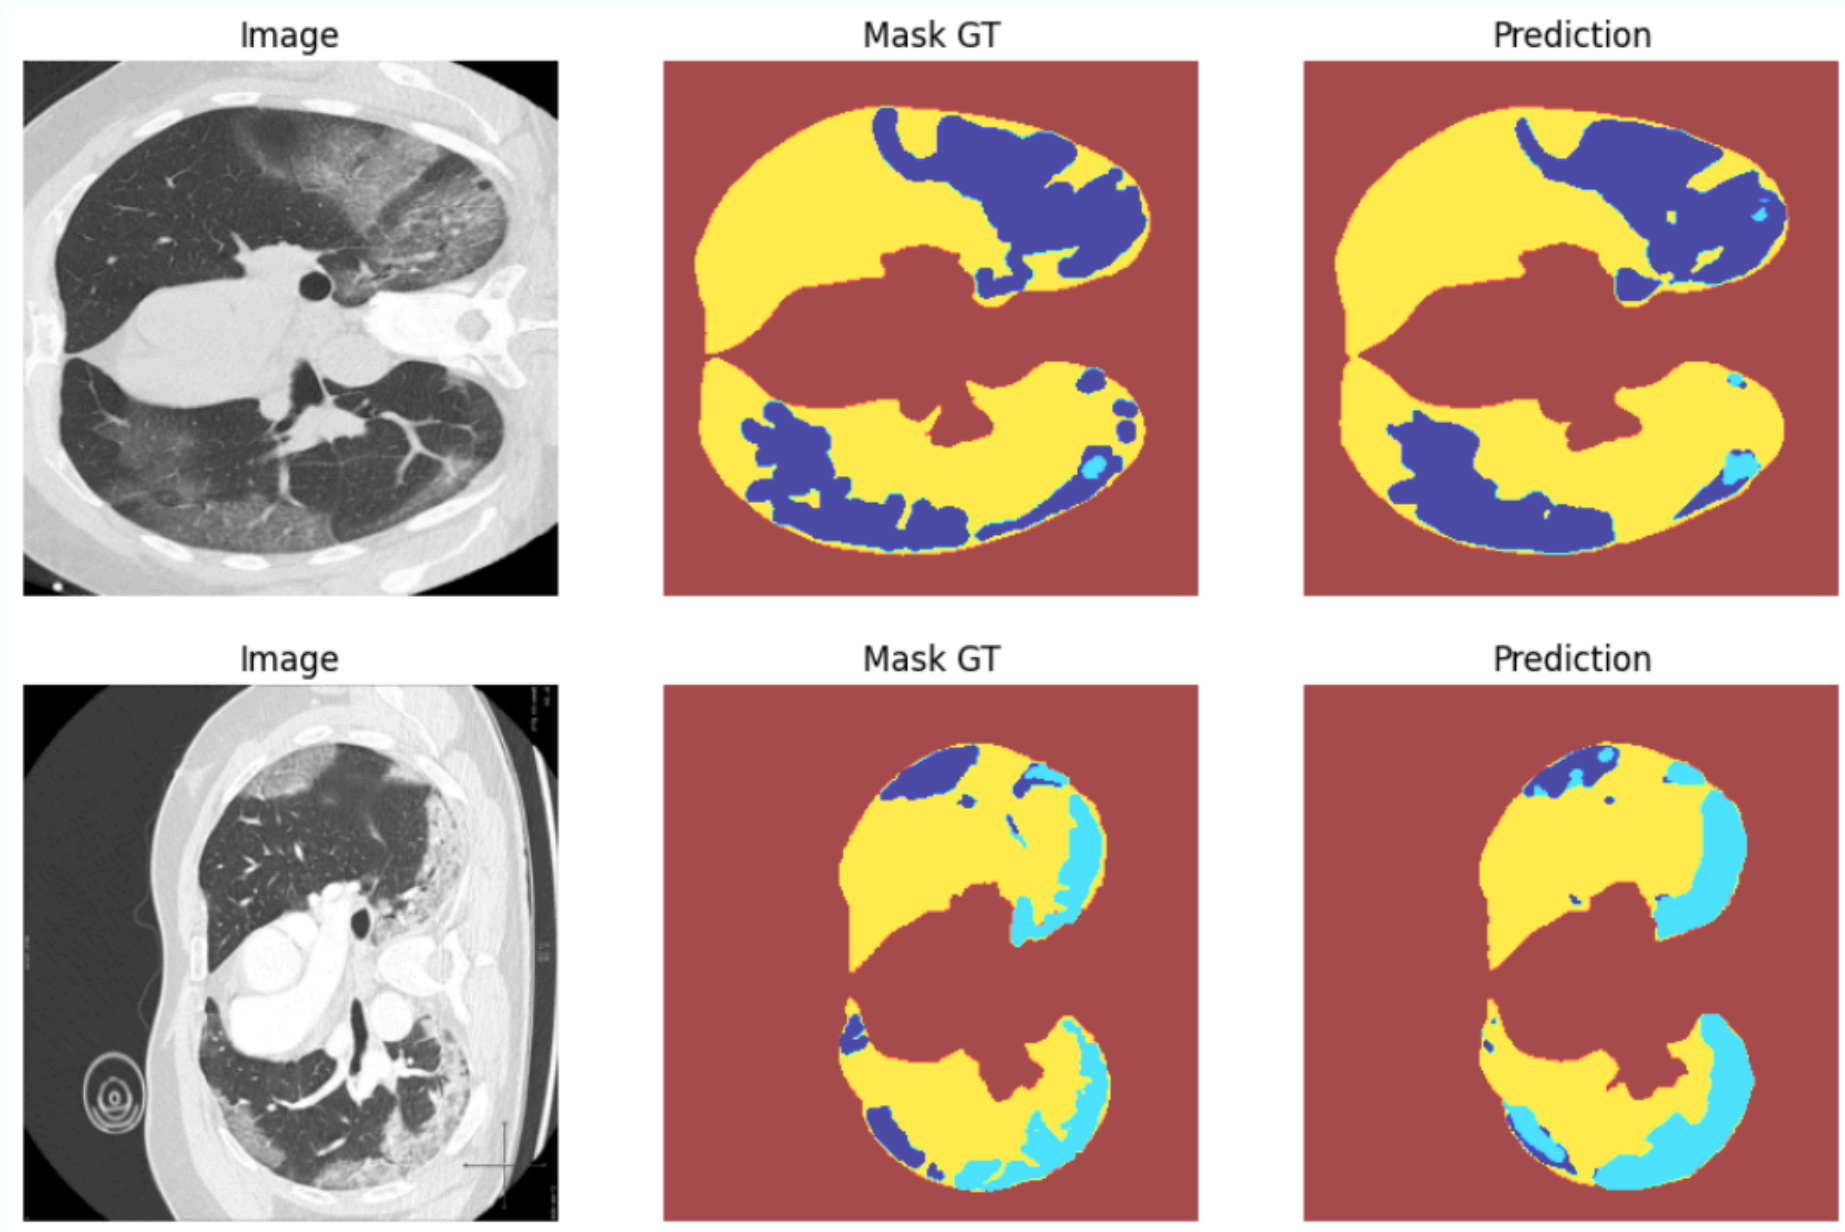
```
        out = model(x)
```

```
        pred = out.argmax(1).cpu()
```

```
        visualize_batch(x.cpu(), y, pred, max_images=8)
```

```
        if b_idx >= 4: # mostrar primeras 5 batches
```

```
            break
```



RESULTADOS

01 SE LOGRÓ UNA
MEJORA
PROGRESIVA EN LA
MÉTRICA DICE A LO
LARGO DE 60
ÉPOCAS.

02 EL MODELO
ALCANZÓ UN DICE
PROMEDIO CERCANO
AL 0.84,
MOSTRANDO BUENA
CAPACIDAD DE
SEGMENTACIÓN EN
LESIONES
PULMONARES.

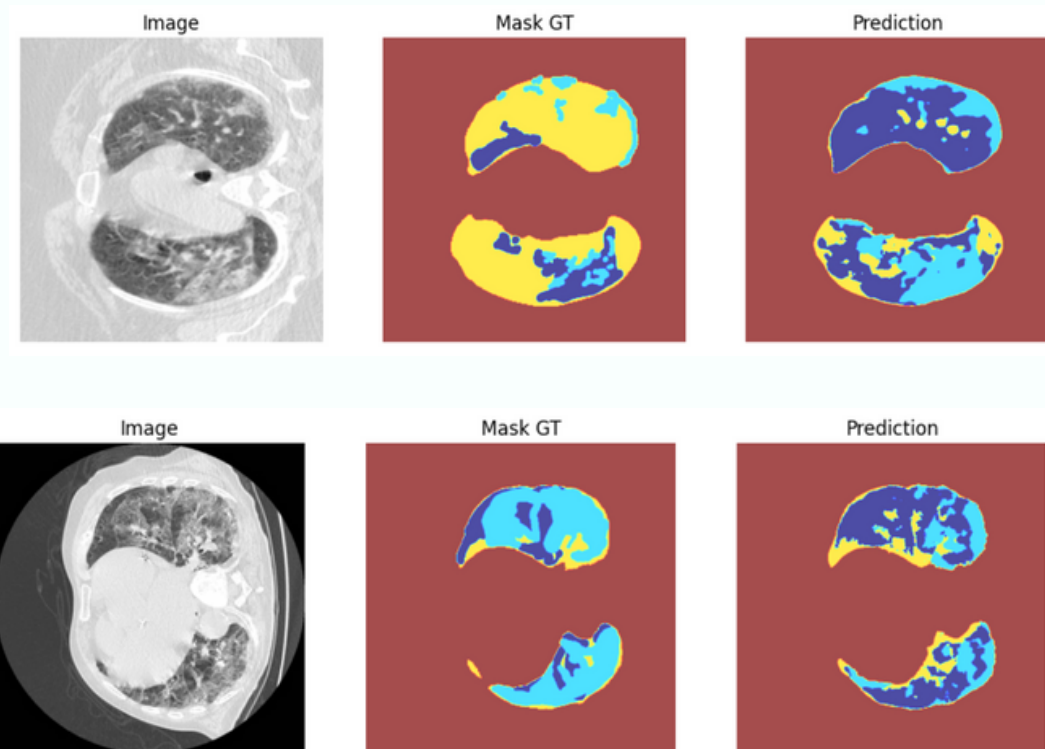


IMPLICACIONES ÉTICAS



MUCHAS
GRACIAS

Limitaciones del modelo



🤖 Why These Outliers Happen

1. Ambiguity in the CT image

- Some slices genuinely look similar across classes (e.g. normal tissue vs. early-stage opacities).
- Even human radiologists sometimes disagree on labeling.

2. Label noise in the ground truth

- Masks in public COVID datasets are often annotated quickly.
- Boundary regions (yellow vs. blue vs. cyan in your case) are **hard to separate**, and your model may be predicting something plausible that just doesn't match the annotation perfectly.

3. Class imbalance still bites

- If one class is underrepresented (say cyan lesions are rare), the model may under- or over-predict it in specific cases.

4. 3D context is missing

- Your UNet++ only sees **one 2D slice at a time**.
- In lungs, some patterns are only clear when you see adjacent slices. Without that, the model may confuse diffuse lesions vs. dense opacities.